



# Introducing Exactly Once Semantics in Apache<sup>®</sup> Kafka<sup>™</sup>

---

Jason Gustafson, Apurva Mehta, Guozhang Wang, and Sriram Subramaniam

Matthias J. Sax | Software Engineer

[matthias@confluent.io](mailto:matthias@confluent.io)

 [@MatthiasJSax](https://twitter.com/MatthiasJSax)

# Outline

---

- Kafka's existing delivery semantics.
- What's new?
- How do you use it?
- Summary.

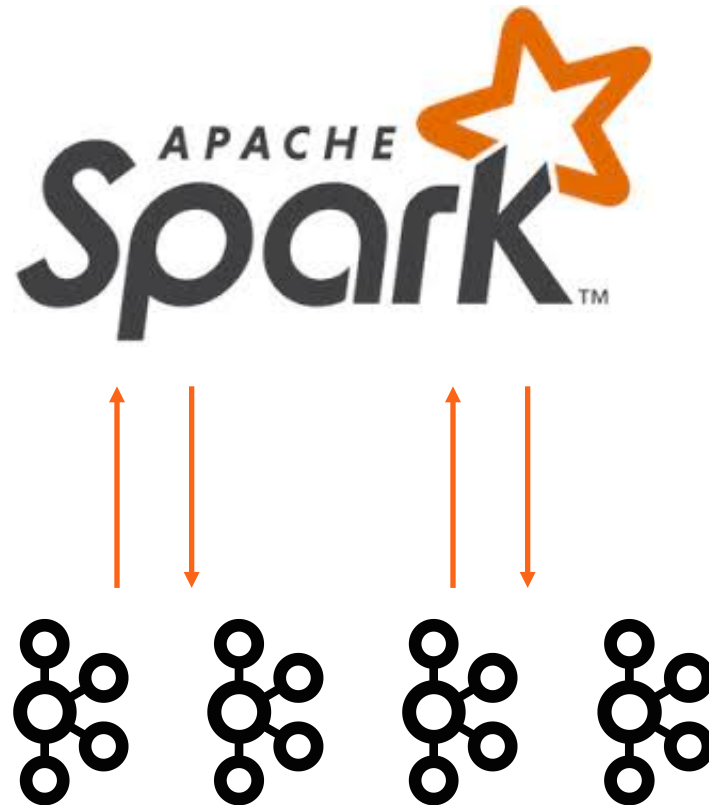
## What Kafka offers today

---

- At-least-once, in-order delivery per partition.
- Producer retries can introduce duplicates.

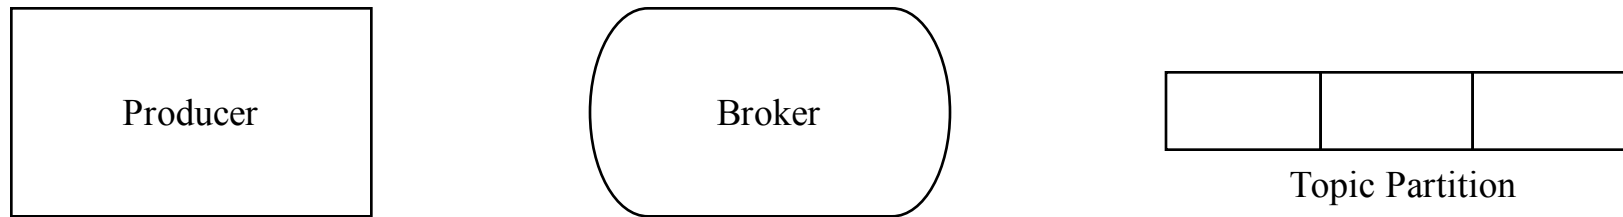
# Stream Processing with Apache Spark

Read-process-write pattern:



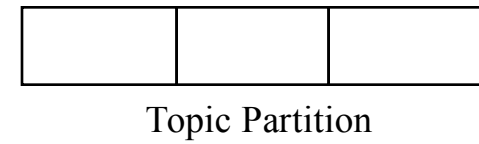
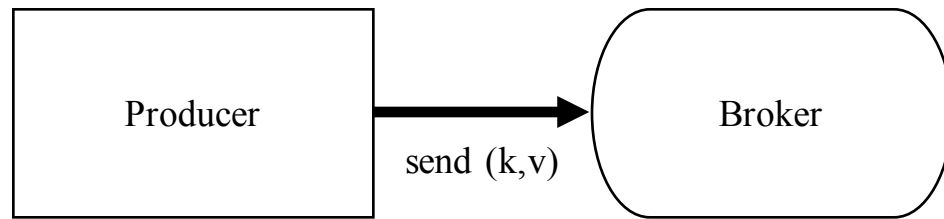
# Example: duplicate write

---



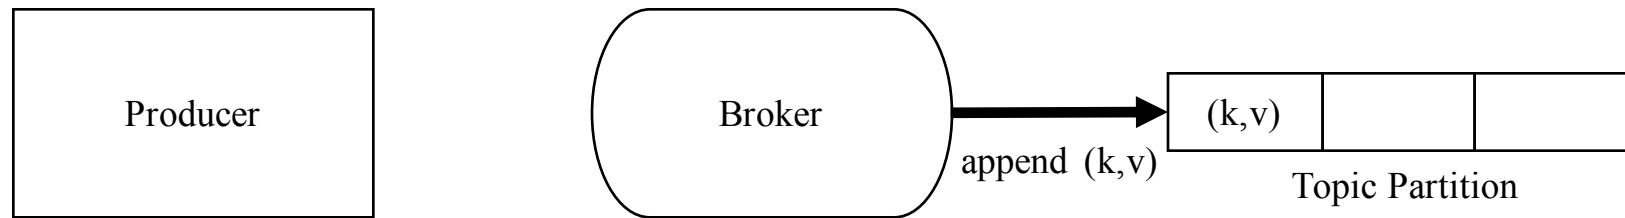
# Example: duplicate write

---





# Example: duplicate write

---



## Example: duplicate write

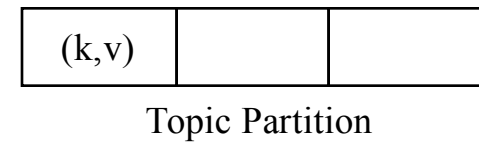
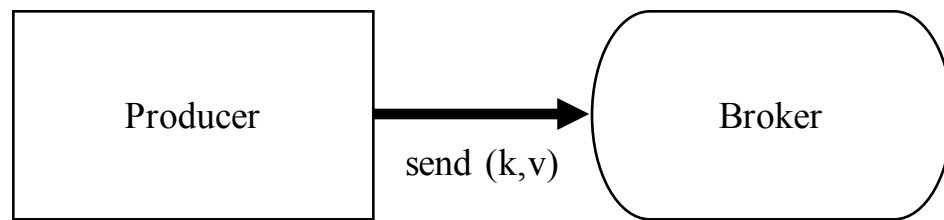
---





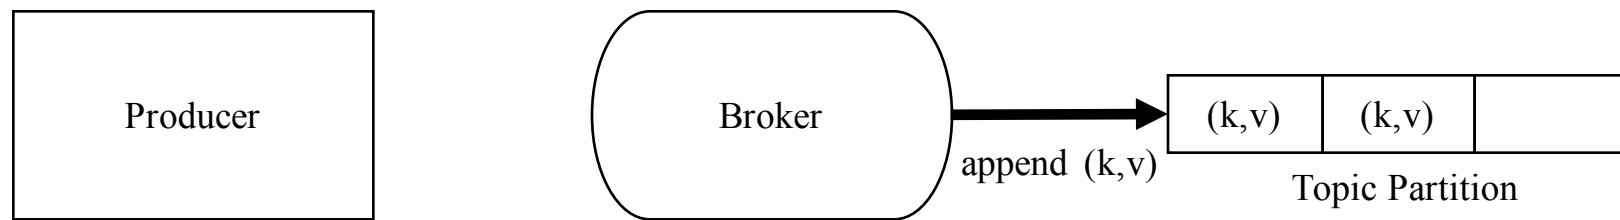
# Example: duplicate write

---



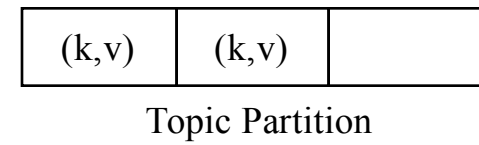
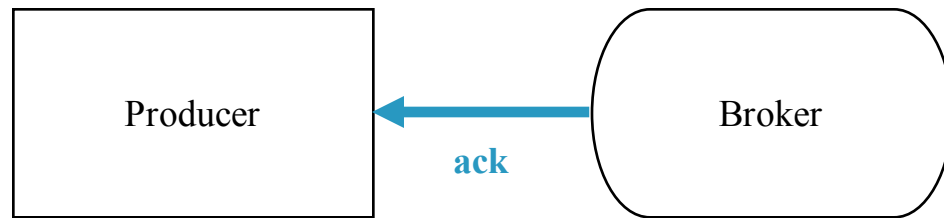
## Example: duplicate write

---



## Example: duplicate write

---



## Why improve?

---

- Stream processing is becoming a bigger part of the data landscape.
- Apache Kafka is the foundation for such stream processing.
- Strengthening Kafka's semantics expands the universe of streaming applications.

---

# What's new?

## What's new

---

- Idempotent producer: exactly-once writes.
- Transactional producer: Atomic writes across multiple partitions.
- Exactly-once stream processing: read-process-write.

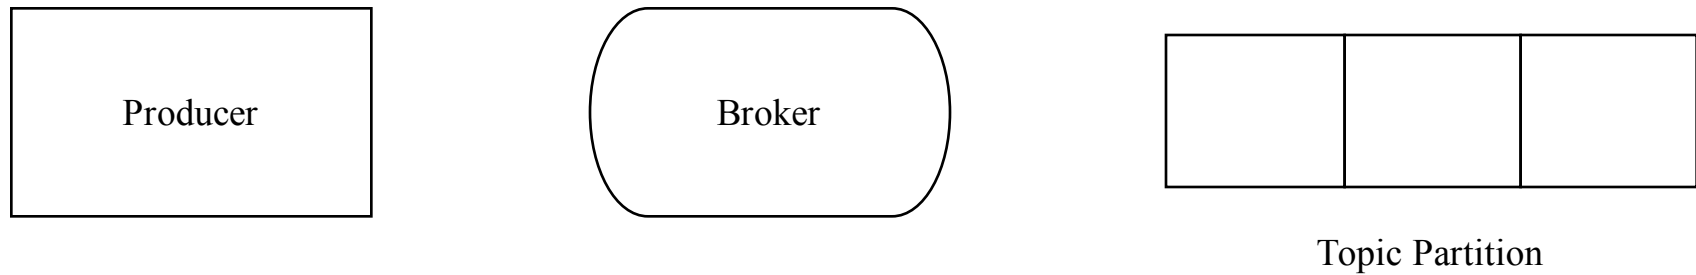
What's new

---

# Idempotent Producer

# Example: Idempotent Producer

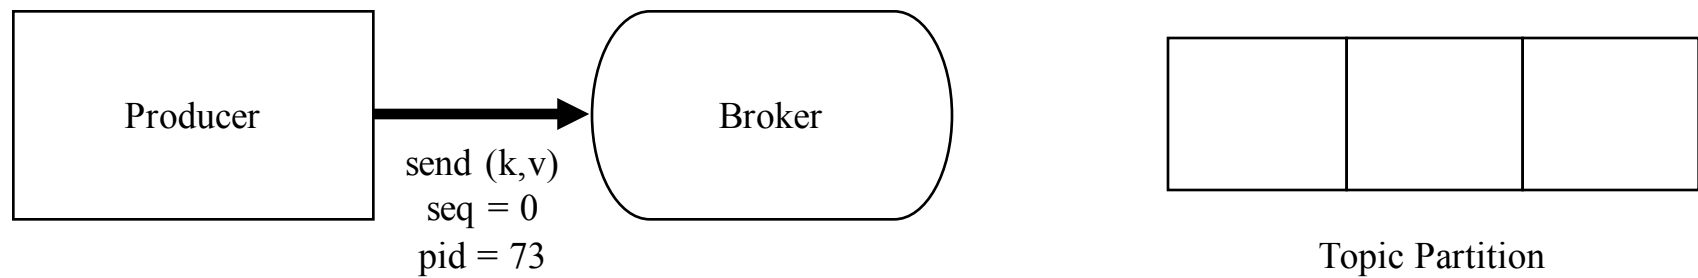
---





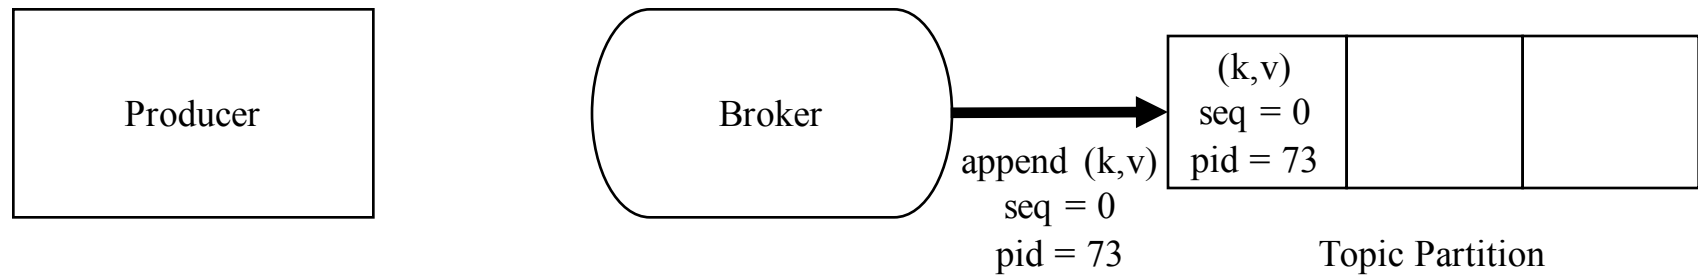
# Example: Idempotent Producer

---

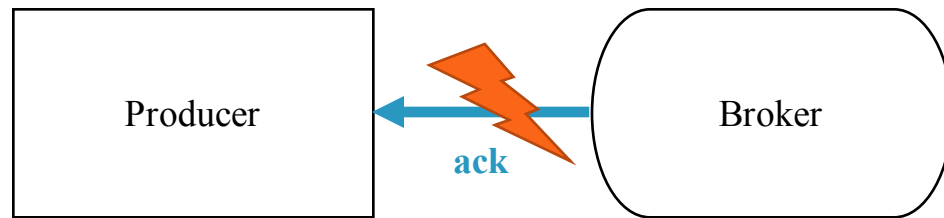


# Example: Idempotent Producer

---



# Example: Idempotent Producer

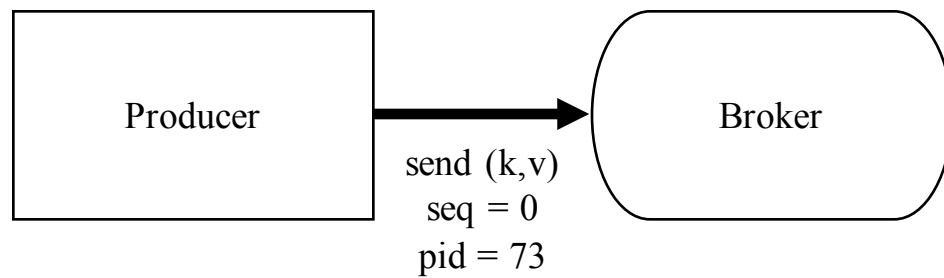


(k,v)		
seq = 0		
pid = 73		

Topic Partition

# Example: Idempotent Producer

---




(k,v)		
seq = 0		
pid = 73		

Topic Partition

# Example: Idempotent Producer

---



(k,v)		
seq = 0		
pid = 73		

Topic Partition

What's new

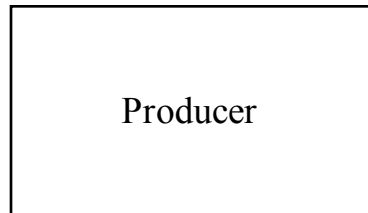
---

# Atomic Multi-Partition Writes

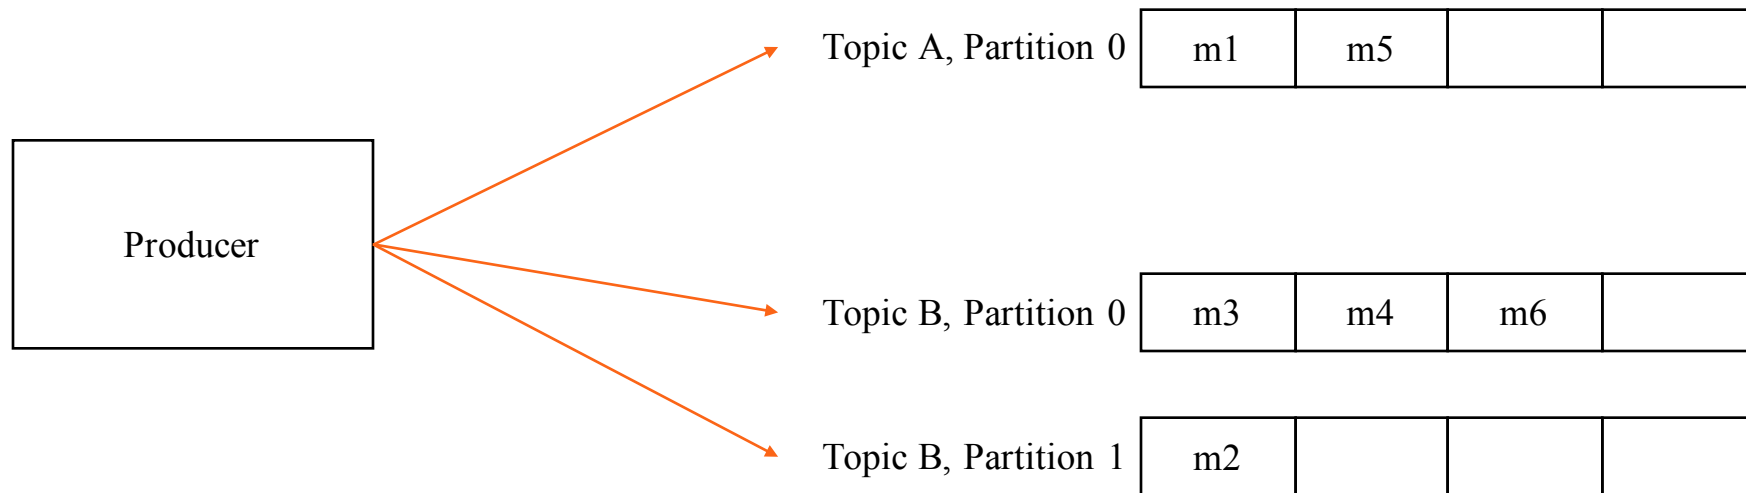
(aka “transactions”)

# Atomic Multi-Partition Writes

---

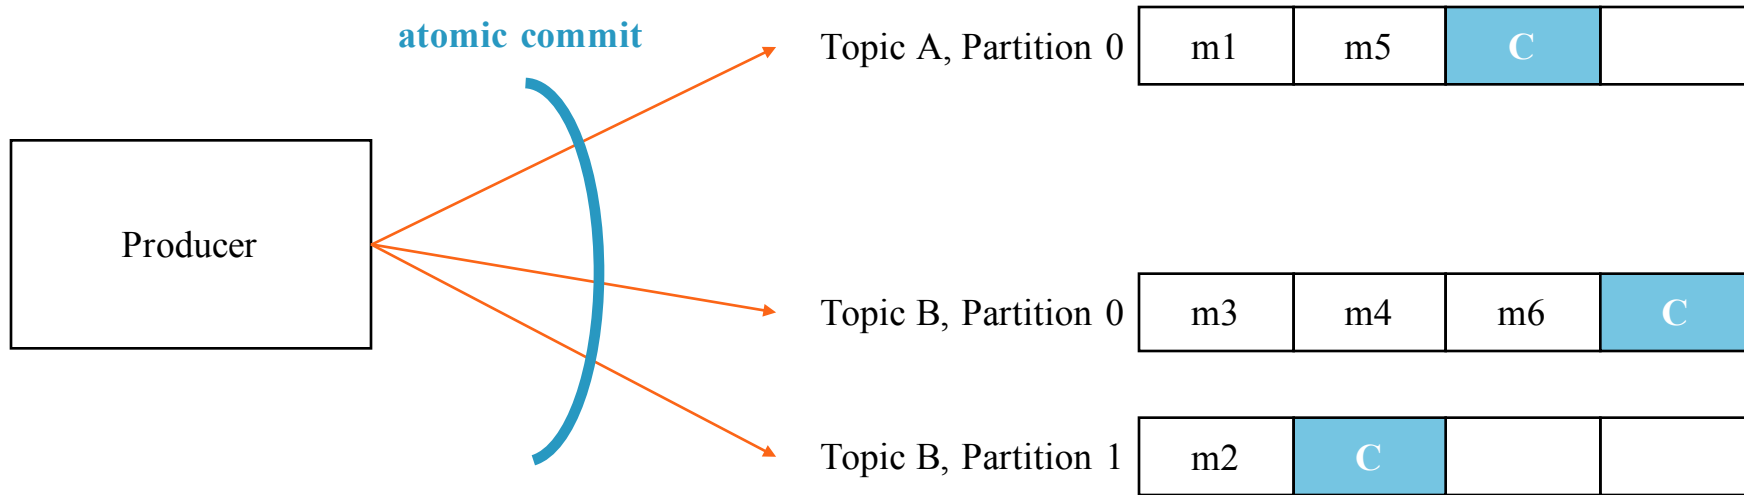


# Atomic Multi-Partition Writes

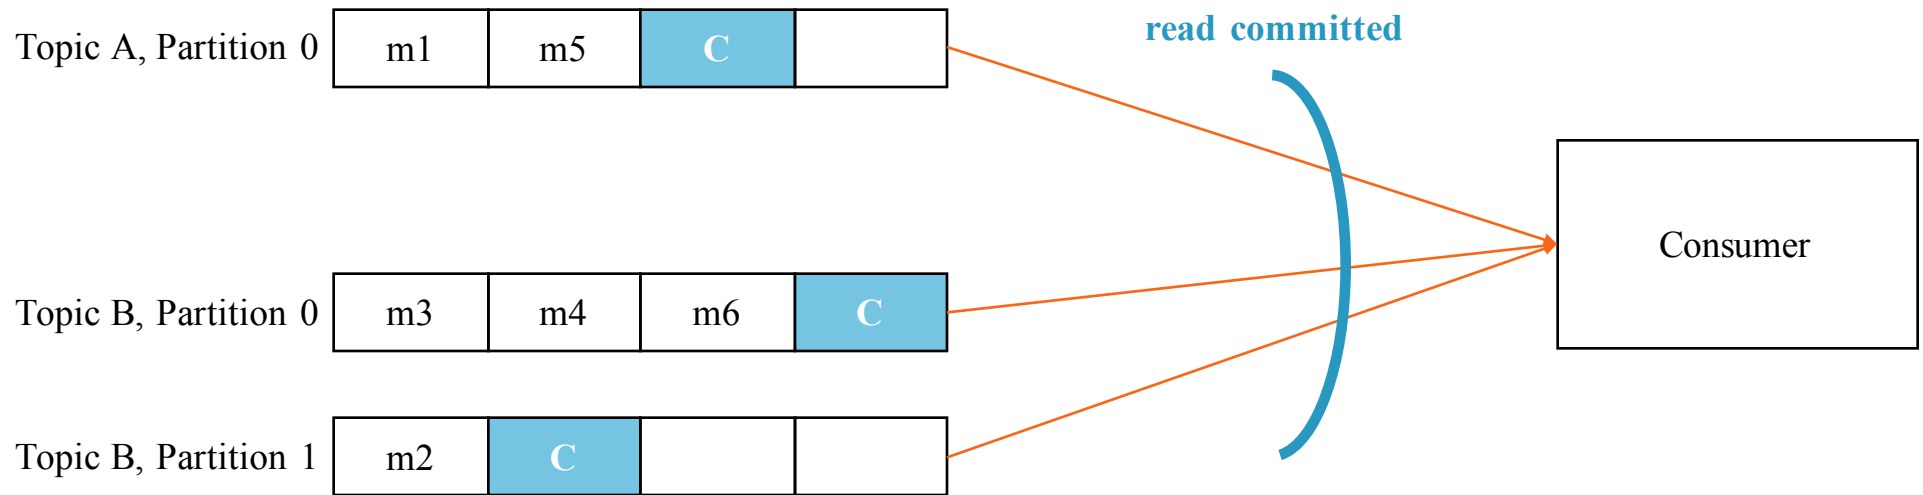




# Atomic Multi-Partition Writes



# Atomic Multi-Partition Writes



# Transactional API

---

```
producer.initTransactions();
try {
    producer.beginTransaction();
    producer.send(record0);
    producer.send(record1);
    producer.sendOffsetsToTxn(...);
    producer.commitTransaction();
} catch (ProducerFencedException e) {
    producer.close();
} catch (KafkaException e) {
    producer.abortTransaction();
}
```

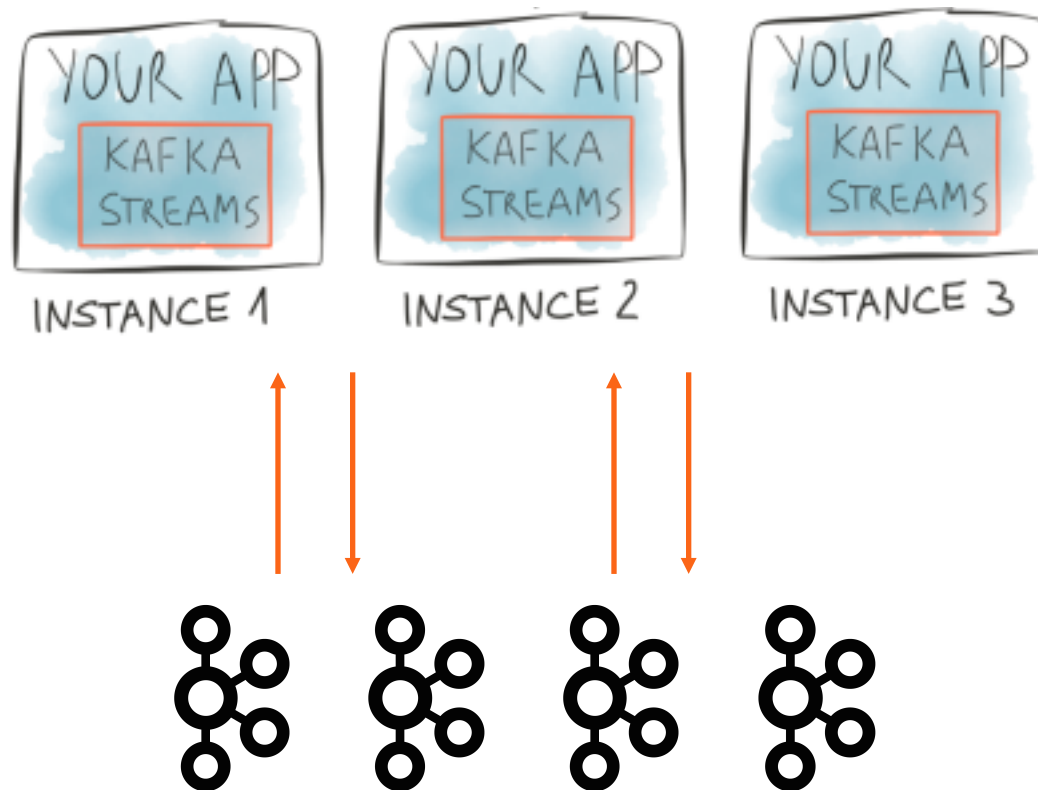
## How to use exactly-once capabilities:

---

- Streams API *(the easiest way to use exactly-once semantics)*
  - Config parameter `processing.mode = "exactly_once"`
- Idempotent Producer
  - Config parameter `enable.idempotence = true`
- Transactional Producer
  - Config parameter `transactional.id = "my-unique-tid"`
  - ***And Transactional API (hard to use!)***
- Transactional Consumer
  - Config parameter `isolation.level = "read_committed"`  
(default: "read\_uncommitted")

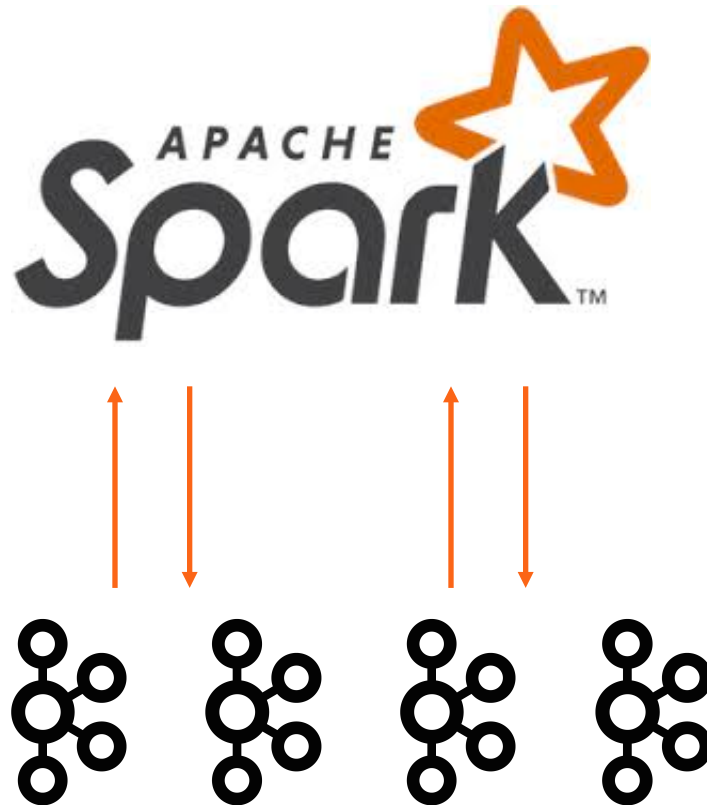
# Stream Processing with Kafka's Streams API

Transactional read-process-write-commit pattern:



# Stream Processing with Apache Spark

Transactional read-process-write-commit pattern:



## When to use this?

---

Available in Kafka 0.11, June 2017. Try it out!

## Putting it together

---

- We understood Kafka's existing delivery semantics.
- Learned how these have been strengthened.
- Learned how the new semantics work.
- Saw, it's easy to use with higher level APIs like Kafka Streams or Apache Spark.





Thank You

---

We are hiring!